




Unit 3: Application Areas

3a: Vision

Perception

- In general we (as humans) perceive a large number of inputs
 - visual, auditory, smell, touch, taste, ESP ;-)
- Somehow we make sense of it all and do the right thing
- Perception is the action of taking raw inputs and producing some beliefs about the world

How images are formed



- Pinhole camera: point P in world shows up as image P'
 - $P=(X,Y,Z)$; $P'=(x,y,f)$
 - $x=-fX/Z$, $y=-fY/Z$
 - The scale and distance is ambiguous

Vision as perception


- We can think of our percepts as arising from some view of the world W

$$P = g(W)$$

- If g is “graphics”, perhaps we can do vision as inverse graphics?

$$W = g^{-1}(P) ?$$

- Problem: Many worlds are possible



Vision as perception


- We can think of our percepts as arising from some view of the world W

$$P = g(W)$$

- If g is “graphics”, perhaps we can do vision as inverse graphics?


$$W = g^{-1}(P) ?$$

- Problem: Many worlds are possible



Vision as perception

- We can think of our percepts as arising from some view of the world W
 $P = g(W)$
- If g is “graphics”, perhaps we can do vision as inverse graphics?
 $W = g^{-1}(P)$?
- Problem: Many worlds are possible



Vision as Bayesian Inference

- Can just treat this as an inference problem
 - $P(W|S) = \propto P(S|W) P(W)$
- The only problem then is trying to figure out priors over world states
 - This can be tricky

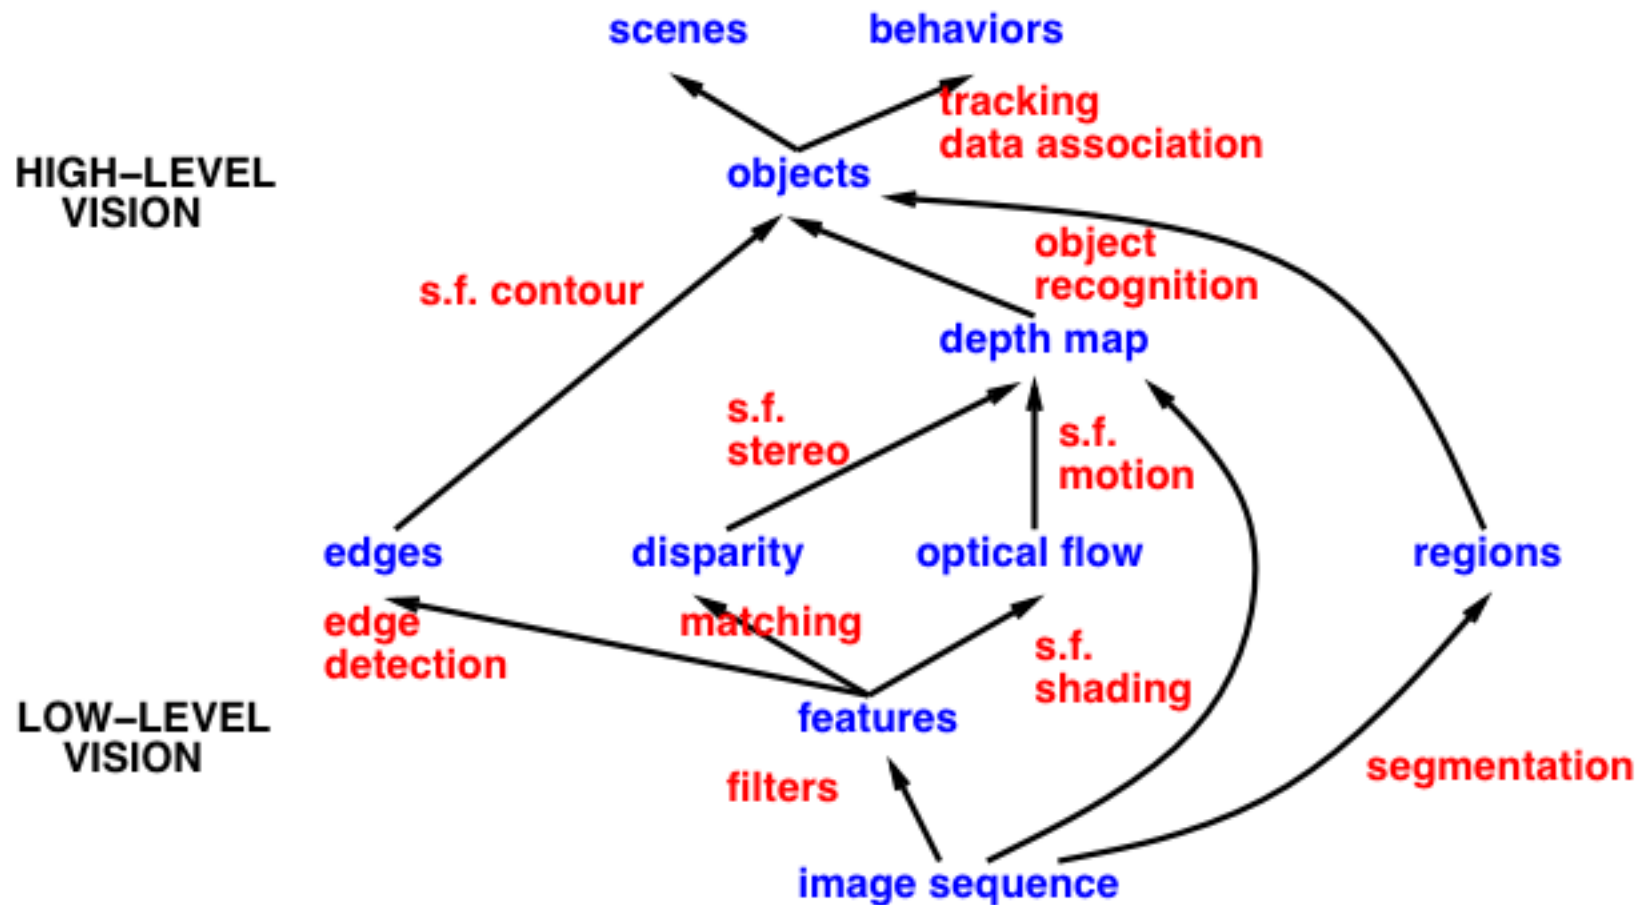
Another approach to vision

- What do you need vision for?
 - Moving around (navigation)
 - Moving things (manipulation)
 - Picking out objects (recognition/identification)
- Don't bother to recover exact scene
- Just get the information needed to do the job
 - Feature extraction, machine learning

What you can get from vision

- Depth (how far away is something)
- Shape (is it flat? curved?)
- Boundaries (edges)
- Objects (putting together information)
- Tracking
- Action recognition

The visual hierarchy




Images




[illegible]


1-dimensional signal





1-dimensional signal





Smoothing it out: averaging





Smoothing it out: averaging



Smoothing it out: averaging



Smoothing it out: averaging



Convolution



- Method of applying a filter to a signal
- Discrete-case equation

$$y(t) = \sum_{n=-\infty}^{\infty} x(t)h(t-n)$$

also written $y = x * h$

Note: convolution is commutative: $y = x * h = h * x$


Number of points if h, x finite: $n_y = n_x + n_h - 1$



Raw image




2-D Gaussian function





$$G(x,y)=G(x)G(y)$$

Gaussian-smoothed image



Detecting edges: 1-D



Detecting edges: 1-D

- Take first derivative of smoothed signal

$$y=(x*g)'$$

- can be estimated by derivative filter $d=[-1 \ 1]$

$$y=(x*g)*d$$

- Theorem: derivative of convolved signal is equivalent to convolving signal with derivative of filter

$$y=x*(g')\approx x*(g*d)$$

- Combine two steps into 1!

Detecting edges: 2-D

- Remember: gaussian components independent


$$G(x,y)=G(x)G(y)$$

- How do you find vertical edges?
 - Derivative in x-direction will change


$$G_x(x,y)=G'(x)G(y)$$

- Similar derivative in y-direction


Gaussian Derivatives



Vertical lines: x-derivative






Horizontal lines: y-derivative




Non-horizontal/vertical lines


- Can combine x and y components to get gradient — normal to edge



- Angle of gradient is 90 degrees off from edge


Combined derivatives





Canny edge detection

- Smooth with gaussian
- Find gradient of image (I.e., derivative in x and y directions)
- Find edge pixels next to each other that point in similar direction



Thresholding

- Sometimes need binary decision on edges
 - Doesn't always work well, might need something more clever





Image segmentation




Cues from prior knowledge

shading





Cues from prior knowledge



Cues from prior knowledge


- Shading: assumes uniform reflectance
- Contour: assumes minimal curvature
- Texture: assumes uniform texture
- Motion: assumes rigid bodies, continuous motion
- Stereo: assumes solid, contiguous, non-repeating bodies

Motion: dx/dt , dy/dt




optical
flow


Stereo vision



Stereo vision: depth perception




Texture



- Patterns repeat in environment
- Use correspondences between parts of pattern to determine orientation of surface

Texture

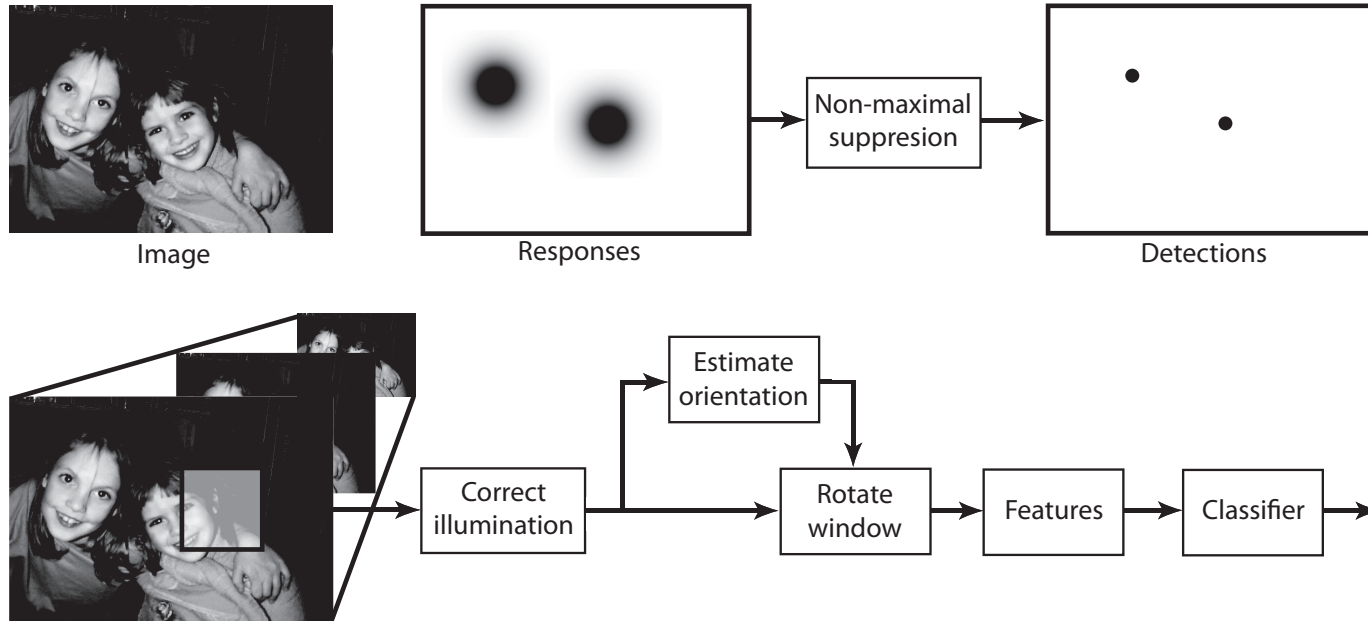


Object recognition

- First: segment the world
- Then: put the pieces back together and figure out what's there
 - Biometric identification: fingerprints, iris scans
 - Content-based image retrieval:
 - Find a photo with “X” in it
 - Handwriting recognition
 - E.g. zip codes


Image content filtering

- How would you find faces in a picture?
- How would you find pictures of trees?




Feature-based recognition

- Extract features (e.g., edges, surfaces, etc)
- Compare against a shape library
- Which of these is not like the other?

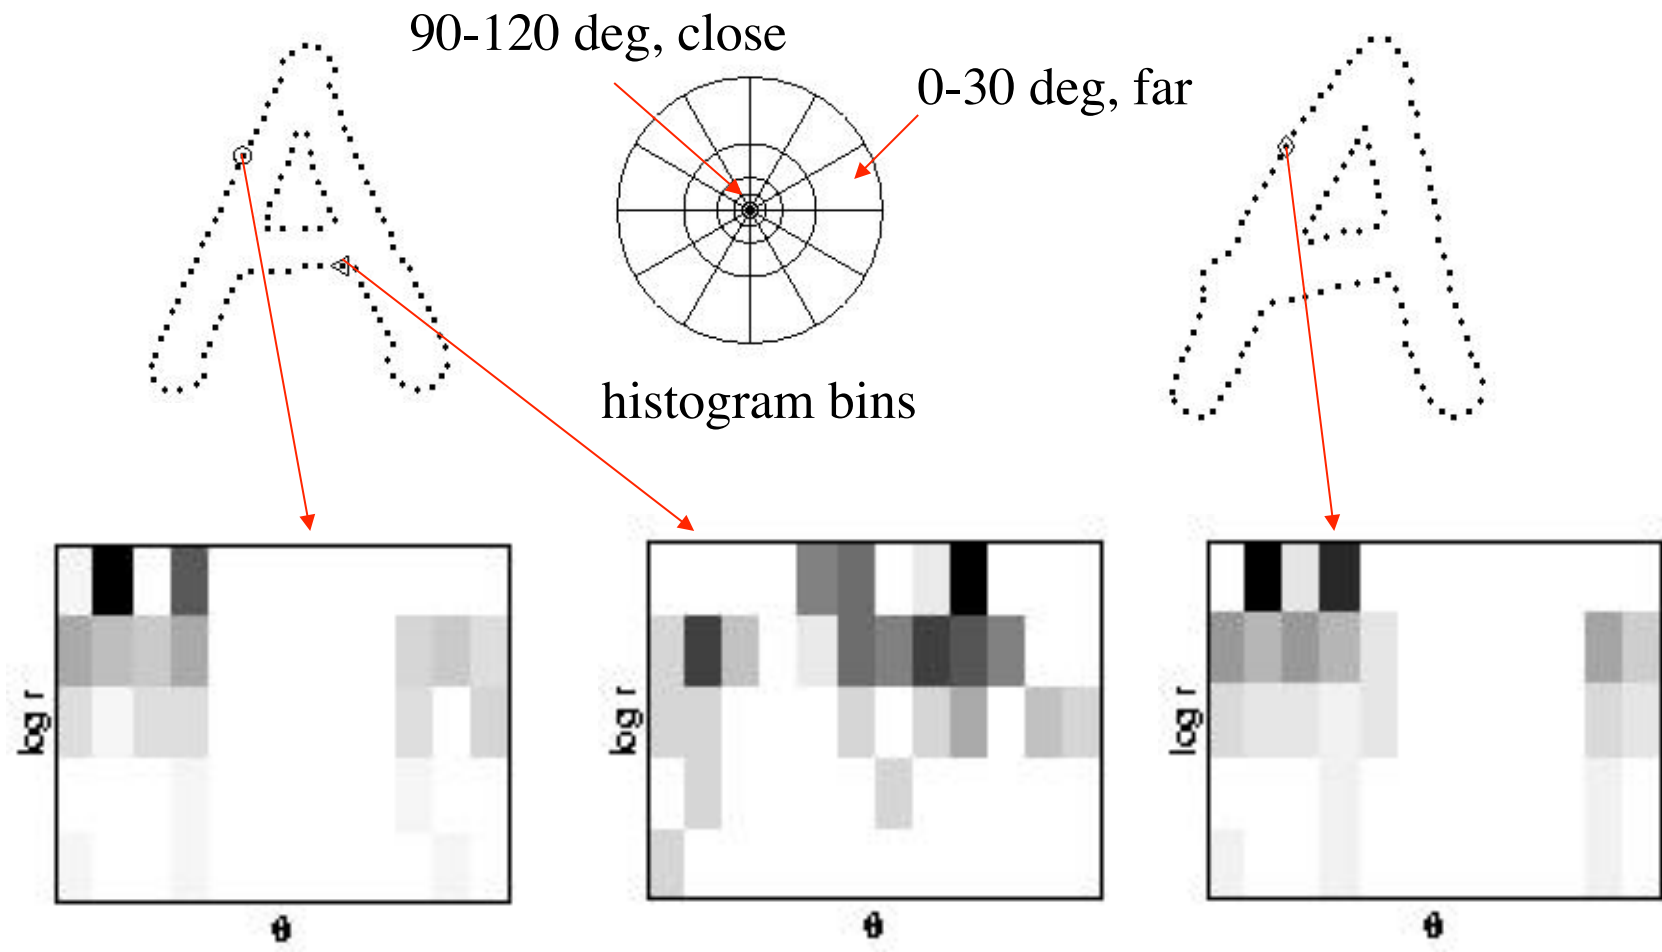


Shape matching

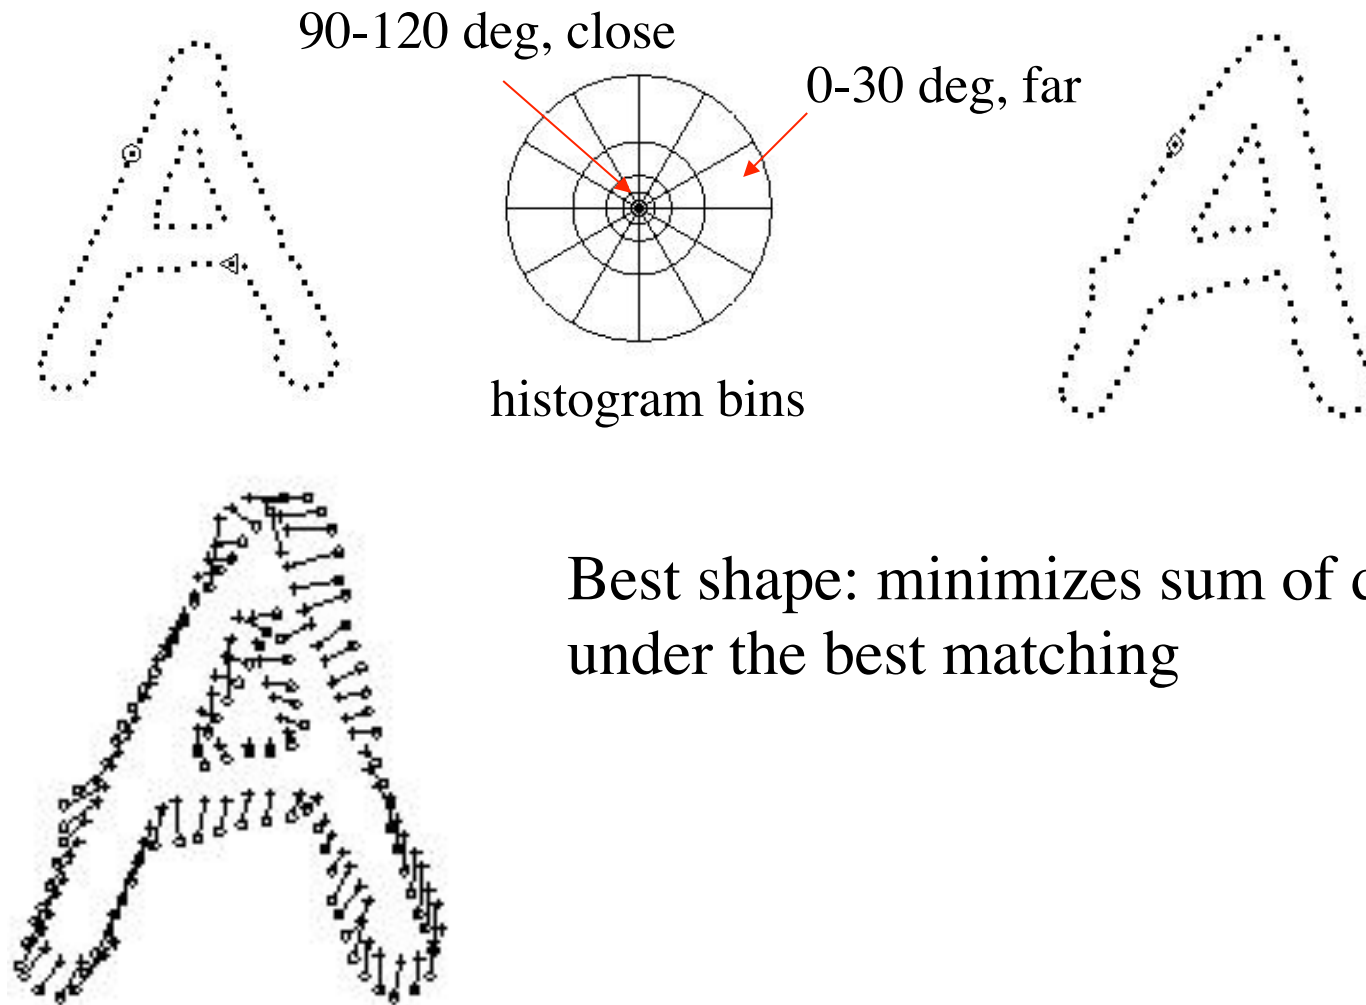


- Don't try to recognize lines
- Distribute a number of points equally around shape
- At every point, take a histogram of the surrounding points
 - Group by angle, log distance


Shape matching



Shape matching




Pose estimation



- Once you match up an object, can you tell which way it's facing
- Determine rotation R and translation t

Computer vision in the real world



What would you need to do this?