

# Multi-Class Logistic Regression and Perceptron

Instructor: Wei Xu

Some slides adapted from Dan Jurfasky, Brendan O'Connor and Marine Carpuat

# MultiClass Classification

- Q: what if we have more than 2 categories?
  - Sentiment: Positive, Negative, Neutral
  - Document topics: Sports, Politics, Business, Entertainment, ...

Q: How to easily do Multi-label classification?

# Two Types of MultiClass Classification

- Multi-label Classification
  - each instance can be assigned more than one labels
- Multinomial Classification
  - each instance appears in exactly one class (classes are exclusive)

# Multinomial Classification

- Pretty straightforward with Naive Bayes.

$$P(\text{spam}|D) \propto P(\text{spam}) \prod_{w \in D} P(w|\text{spam})$$

# Log-Linear Models

$$P(y|x) \propto e^{w \cdot f(x,y)}$$

$$P(y|x) = \frac{1}{Z(w)} e^{w \cdot f(x,y)}$$

# Multinomial Logistic Regression

$$P(y|x) \propto e^{w \cdot f(x,y)}$$

$$P(y|x) = \frac{1}{Z(w)} e^{w \cdot f(x,y)}$$

$$P(y|x) = \frac{e^{w \cdot f(x,y)}}{\sum_{y' \in Y} e^{w \cdot f(x,y')}} \quad \text{normalization term (Z) so that probabilities sum to 1}$$

# (a.k.a) Softmax Regression



WIKIPEDIA  
The Free Encyclopedia

[Main page](#)  
[Contents](#)  
[Featured content](#)  
[Current events](#)  
[Random article](#)  
[Donate to Wikipedia](#)  
[Wikipedia store](#)

Not logged in [Talk](#) [Contributions](#) [Create account](#) [Log in](#)

Article

[Talk](#)

Read

[Edit](#)

[View history](#)



## Softmax function

From Wikipedia, the free encyclopedia

In [mathematics](#), the **softmax function**, or **normalized exponential function**,<sup>[1]:198</sup> is a generalization of the [logistic function](#) that "squashes" a  $K$ -dimensional vector  $\mathbf{z}$  of arbitrary real values to a  $K$ -dimensional vector  $\sigma(\mathbf{z})$  of real values in the range  $(0, 1)$  that add up to 1. The function is given by

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$

Q: what if there are only 2 categories?

$$P(y = j | x_i) = \frac{e^{w_j \cdot x_i}}{\sum_k e^{w_k \cdot x_i}}$$



Q: what if there are only 2 categories?

$$P(y = 1|x) = \frac{e^{w_1 \cdot x}}{e^{w_0 \cdot x + w_1 \cdot x - w_1 \cdot x} + e^{w_1 \cdot x}}$$

Q: what if there are only 2 categories?

$$P(y = 1|x) = \frac{e^{w_1 \cdot x}}{e^{w_0 \cdot x - w_1 \cdot x} e^{w_1 \cdot x} + e^{w_1 \cdot x}}$$

Q: what if there are only 2 categories?

$$P(y = 1|x) = \frac{e^{w_1 \cdot x}}{e^{w_1 \cdot x} (e^{w_0 \cdot x - w_1 \cdot x} + 1)}$$

Q: what if there are only 2 categories?

$$P(y = 1|x) = \frac{1}{e^{w_0 \cdot x - w_1 \cdot x} + 1}$$

Q: what if there are only 2 categories?

$$P(y = 1|x) = \frac{1}{e^{-w' \cdot x} + 1}$$



Sigmoid (logistic) function

# Multinomial Logistic Regression

- Binary (two classes):
  - We have one feature vector that matches the size of the vocabulary
- Multi-class in practice:
  - one weight vector for each category

$w_{\text{pos}}$

$w_{\text{neg}}$

$w_{\text{neut}}$

In practice, can represent this with one giant weight vector and repeated features for each category.

# Maximum Likelihood Estimation

$$w_{\text{MLE}} = \operatorname{argmax}_w \log P(y_1, \dots, y_n | x_1, \dots, x_n; w)$$

$$= \operatorname{argmax}_w \sum_i \log P(y_i | x_i; w)$$

$$= \operatorname{argmax}_w \sum_i \log \frac{e^{w \cdot f(x_i, y_i)}}{\sum_{y' \in Y} e^{w \cdot f(x_i, y')}}$$

# Multiclass LR Gradient

$$\frac{\partial \mathcal{L}}{\partial w_j} = \sum_{i=1}^D f_j(y_i, d_i) - \sum_{i=1}^D \sum_{y \in Y} f_j(y, d_i) P(y|d_i)$$

empirical feature count

expected feature count



# (a.k.a) Maximum Entropy Classifier

- or MaxEnt
- Math proof of “LR=MaxEnt”:
  - [Klein and Manning 2003]
  - [Mount 2011]

# Perceptron Algorithm


- Very similar to logistic regression
- Not exactly computing gradient



[Rosenblatt 1957]

# Perceptron Algorithm

- Very similar to logistic regression
- Not exactly computing gradient (simpler)



$$P(y=1|x)=1 \text{ if } w \cdot \phi(x) \geq 0$$
$$P(y=1|x)=0 \text{ if } w \cdot \phi(x) < 0$$

$$P(y=1|x) = \frac{e^{w \cdot \phi(x)}}{1 + e^{w \cdot \phi(x)}}$$

# Perceptron vs. LR

- The Perceptron is an online learning algorithm.
- Standard Logistic Regression is not

# Online Learning

- The Perceptron is an online learning algorithm.
- Logistic Regression is not:

this update is effectively the same as “ $w \leftarrow w + y_i x_i$ ”

$$\begin{aligned} w_{\text{MLE}} &= \operatorname{argmax}_w \log P(y_1, \dots, y_d | x_1, \dots, x_d; w) \\ &= \operatorname{argmax}_w \sum_i y_i \log p_i + (1 - y_i) \log(1 - p_i) \end{aligned}$$

# (Full) Batch Learning

- update parameters after each pass of training set

Initialize weight vector  $w = 0$

Create features

Loop for  $K$  iterations

    Loop for all training examples  $x_i, y_i$

        ...

        update\_weights( $w$ )

# Online Learning

- update parameters for each training example

Initialize weight vector  $w = 0$

Create features

Loop for  $K$  iterations

    Loop for all training examples  $x_i, y_i$

        ...

        update\_weights( $w, x_i, y_i$ )

# Perceptron Algorithm

- Very similar to logistic regression
- Not exactly computing gradient

The diagram illustrates the Perceptron update rule:  $w \leftarrow w + y \varphi(x)$ . It includes three callout boxes: one pointing to the  $w$  on the right with the text "weights", one pointing to the  $y$  with the text "label", and one pointing to the  $\varphi(x)$  term with the text "features of a training example x".

$$w \leftarrow w + y \varphi(x)$$

weights

label

features of a training example x

If  $y = 1$ , increase the weights for features in  $\varphi(x)$

If  $y = -1$ , decrease the weights for features in  $\varphi(x)$




# Perceptron Algorithm

- Very similar to logistic regression
- Not exactly computing gradient

Initialize weight vector  $w = 0$

Loop for  $K$  iterations

Loop For all training examples  $x_i$

if  $\text{sign}(w * x_i) \neq y_i$  

$w += y_i * x_i$

# The Intuition

- For a given example, makes a prediction, then checks to see if this prediction is correct.
- If the prediction is correct, do nothing.
- If the prediction is incorrect, change its parameters so that it would do better on this example next time around.

# Perceptron (vs. LR)


- Only hyperparameter is maximum number of iterations (LR also needs learning rate)

# Perceptron (vs. LR)


- Only hyperparameter is maximum number of iterations (LR also needs learning rate)
- Guaranteed to converge if the data is linearly separable (LR always converge)

# Linear Separability


i




ii



iii



iv



# What does “converge” mean?

- It means that it can make an entire pass through the training data without making any more updates.
- In other words, it has correctly classified every training example.
- Geometrically, this means that it was found some hyperplane that correctly segregates the data into positive and negative examples

# What if non linearly separable?

- In real-world problem, this is nearly always the case.
- The perceptron will not be able to converge.

Q: Then, when to stop?