

Log-linear Models (Review)

Instructor: Wei Xu

Many slides from Michael Collins

Overview

- ▶ Log-linear models
- ▶ Parameter estimation in log-linear models
- ▶ Smoothing/regularization in log-linear models

The General Problem

- ▶ We have some **input domain** \mathcal{X}
- ▶ Have a finite **label set** \mathcal{Y}
- ▶ Aim is to provide a **conditional probability** $p(y \mid x)$
for any x, y where $x \in \mathcal{X}, y \in \mathcal{Y}$

Feature Vector Representations

- ▶ Aim is to provide a conditional probability $p(y \mid x)$ for “decision” y given “history” x
- ▶ A **feature** is a function $f_k(x, y) \in \mathbb{R}$
(Often **binary features** or **indicator functions** $f_k(x, y) \in \{0, 1\}$).
- ▶ Say we have m features f_k for $k = 1 \dots m$
 \Rightarrow A **feature vector** $f(x, y) \in \mathbb{R}^m$ for any x, y

features are a property of both observation x and the candidate output class y

Parameter Vectors

- ▶ Given features $f_k(x, y)$ for $k = 1 \dots m$, also define a **parameter vector** $v \in \mathbb{R}^m$
- ▶ Each (x, y) pair is then mapped to a “score”

all possible m-dimensional
real value vectors

$$v \cdot f(x, y) = \sum_k v_k f_k(x, y)$$

However, this doesn't
produce a legal probability

Log-Linear Models

- ▶ We have some input domain \mathcal{X} , and a finite label set \mathcal{Y} . Aim is to provide a conditional probability $p(y \mid x)$ for any $x \in \mathcal{X}$ and $y \in \mathcal{Y}$.
- ▶ A feature is a function $f : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$
(Often binary features or indicator functions $f_k : \mathcal{X} \times \mathcal{Y} \rightarrow \{0, 1\}$).
- ▶ Say we have m features f_k for $k = 1 \dots m$
 \Rightarrow A feature vector $f(x, y) \in \mathbb{R}^m$ for any $x \in \mathcal{X}$ and $y \in \mathcal{Y}$.
- ▶ We also have a **parameter vector** $v \in \mathbb{R}^m$
- ▶ We define

$$p(y \mid x; v) = \frac{e^{v \cdot f(x, y)}}{\sum_{y' \in \mathcal{Y}} e^{v \cdot f(x, y')}}$$

Softmax!

Exercise

Why the name?

$$\log p(y \mid x; v) = \underbrace{v \cdot f(x, y)}_{\text{Linear term}} - \underbrace{\log \sum_{y' \in \mathcal{Y}} e^{v \cdot f(x, y')}}_{\text{Normalization term}}$$

Overview

- ▶ Log-linear models
- ▶ Parameter estimation in log-linear models
- ▶ Smoothing/regularization in log-linear models

Maximum-Likelihood Estimation

- ▶ Maximum-likelihood estimates given training sample $(x^{(i)}, y^{(i)})$ for $i = 1 \dots n$, each $(x^{(i)}, y^{(i)}) \in \mathcal{X} \times \mathcal{Y}$:

$$v_{ML} = \operatorname{argmax}_{v \in \mathbb{R}^m} L(v)$$

where

$$L(v) = \sum_{i=1}^n \log p(y^{(i)} \mid x^{(i)}; v) = \sum_{i=1}^n v \cdot f(x^{(i)}, y^{(i)}) - \sum_{i=1}^n \log \sum_{y' \in \mathcal{Y}} e^{v \cdot f(x^{(i)}, y')}$$

concave function!


Calculating the Maximum-Likelihood Estimates

- Need to maximize:

$$L(v) = \sum_{i=1}^n v \cdot f(x^{(i)}, y^{(i)}) - \sum_{i=1}^n \log \sum_{y' \in \mathcal{Y}} e^{v \cdot f(x^{(i)}, y')}$$

- Calculating gradients:

$$\begin{aligned} \frac{dL(v)}{dv_k} &= \sum_{i=1}^n f_k(x^{(i)}, y^{(i)}) - \sum_{i=1}^n \frac{\sum_{y' \in \mathcal{Y}} f_k(x^{(i)}, y') e^{v \cdot f(x^{(i)}, y')}}{\sum_{z' \in \mathcal{Y}} e^{v \cdot f(x^{(i)}, z')}} \\ &= \sum_{i=1}^n f_k(x^{(i)}, y^{(i)}) - \sum_{i=1}^n \sum_{y' \in \mathcal{Y}} f_k(x^{(i)}, y') \frac{e^{v \cdot f(x^{(i)}, y')}}{\sum_{z' \in \mathcal{Y}} e^{v \cdot f(x^{(i)}, z')}} \\ &= \underbrace{\sum_{i=1}^n f_k(x^{(i)}, y^{(i)})}_{\text{Empirical counts}} - \underbrace{\sum_{i=1}^n \sum_{y' \in \mathcal{Y}} f_k(x^{(i)}, y') p(y' \mid x^{(i)}; v)}_{\text{Expected counts}} \end{aligned}$$



Gradient Ascent Methods

- ▶ Need to maximize $L(v)$ where

$$\frac{dL(v)}{dv} = \sum_{i=1}^n f(x^{(i)}, y^{(i)}) - \sum_{i=1}^n \sum_{y' \in \mathcal{Y}} f(x^{(i)}, y') p(y' | x^{(i)}; v)$$

Initialization: $v = 0$

Iterate until convergence:

- ▶ Calculate $\Delta = \frac{dL(v)}{dv}$
- ▶ Calculate $\beta_* = \operatorname{argmax}_{\beta} L(v + \beta\Delta)$ (Line Search)
- ▶ Set $v \leftarrow v + \beta_*\Delta$


Conjugate Gradient Methods

- ▶ (Vanilla) gradient ascent can be very slow
- ▶ Conjugate gradient methods require calculation of gradient at each iteration, but do a line search in a **direction which is a function of the current gradient, and the previous step taken**.
- ▶ Conjugate gradient packages are widely available
In general: they require a function

$$\text{calc_gradient}(v) \rightarrow \left(L(v), \frac{dL(v)}{dv} \right)$$

and that's about it!

e.g. LBFGS Algorithm
(Limited-memory Broyden-Fletcher-Goldfarb-Shanno)



Overview

- ▶ Log-linear models
- ▶ Parameter estimation in log-linear models
- ▶ Smoothing/regularization in log-linear models

Smoothing in Log-Linear Models

- ▶ Say we have a feature:

$$f_{100}(x, y) = \begin{cases} 1 & \text{if current word } w_i \text{ is base and } y = \text{Vt} \\ 0 & \text{otherwise} \end{cases}$$

- ▶ In training data, base is seen 3 times, with Vt every time
- ▶ Maximum likelihood solution satisfies

$$\sum_i f_{100}(x^{(i)}, y^{(i)}) = \sum_i \sum_y p(y \mid x^{(i)}; v) f_{100}(x^{(i)}, y)$$

- $\Rightarrow p(\text{Vt} \mid x^{(i)}; v) = 1$ for any history $x^{(i)}$ where $w_i = \text{base}$
- $\Rightarrow v_{100} \rightarrow \infty$ at maximum-likelihood solution (most likely)
- $\Rightarrow p(\text{Vt} \mid x; v) = 1$ for any test data history x where $w = \text{base}$

Regularization

- ▶ Modified loss function

$$L(v) = \sum_{i=1}^n v \cdot f(x^{(i)}, y^{(i)}) - \sum_{i=1}^n \log \sum_{y' \in \mathcal{Y}} e^{v \cdot f(x^{(i)}, y')} - \frac{\lambda}{2} \sum_{k=1}^m v_k^2$$

- ▶ Calculating gradients:

$$\frac{dL(v)}{dv_k} = \underbrace{\sum_{i=1}^n f_k(x^{(i)}, y^{(i)})}_{\text{Empirical counts}} - \underbrace{\sum_{i=1}^n \sum_{y' \in \mathcal{Y}} f_k(x^{(i)}, y') p(y' \mid x^{(i)}; v)}_{\text{Expected counts}} - \lambda v_k$$

- ▶ Can run conjugate gradient methods as before
- ▶ Adds a penalty for large weights