

# Maximum Entropy Markov Models (log-linear model for tagging)

Instructor: Wei Xu

Many slides from Michael Collins

# Part-of-Speech Tagging

## INPUT:

Profits soared at Boeing Co., easily topping forecasts on Wall Street, as their CEO Alan Mulally announced first quarter results.

## OUTPUT:

Profits/**N** soared/**V** at/**P** Boeing/**N** Co./**N** ,/, easily/**ADV** topping/**V**  
forecasts/**N** on/**P** Wall/**N** Street/**N** ,/, as/**P** their/**POSS** CEO/**N**  
Alan/**N** Mulally/**N** announced/**V** first/**ADJ** quarter/**N** results/**N** ./.

**N** = Noun

**V** = Verb

**P** = Preposition

**Adv** = Adverb

**Adj** = Adjective

...

# Named Entity Recognition

**INPUT:** Profits soared at Boeing Co., easily topping forecasts on Wall Street, as their CEO Alan Mulally announced first quarter results.

**OUTPUT:** Profits soared at [Company Boeing Co.], easily topping forecasts on [Location Wall Street], as their CEO [Person Alan Mulally] announced first quarter results.

# Named Entity Extraction as Tagging

## INPUT:

Profits soared at Boeing Co., easily topping forecasts on Wall Street, as their CEO Alan Mulally announced first quarter results.

## OUTPUT:

Profits/NA soared/NA at/NA Boeing/SC Co./CC ,/NA easily/NA  
topping/NA forecasts/NA on/NA Wall/SL Street/CL ,/NA as/NA  
their/NA CEO/NA Alan/SP Mulally/CP announced/NA first/NA  
quarter/NA results/NA ./NA

- NA = No entity
- SC = Start Company
- CC = Continue Company
- SL = Start Location
- CL = Continue Location

# Our Goal

## Training set:

1 Pierre/**NNP** Vinken/**NNP** ,/, 61/**CD** years/**NNS** old/**JJ** ,/, will/**MD** join/**VB** the/**DT** board/**NN** as/**IN** a/**DT** nonexecutive/**JJ** director/**NN** Nov./**NNP** 29/**CD** ./.

2 Mr./**NNP** Vinken/**NNP** is/**VBZ** chairman/**NN** of/**IN** Elsevier/**NNP** N.V./**NNP** ,/, the/**DT** Dutch/**NNP** publishing/**VBG** group/**NN** ./.

3 Rudolph/**NNP** Agnew/**NNP** ,/, 55/**CD** years/**NNS** old/**JJ** and/**CC** chairman/**NN** of/**IN** Consolidated/**NNP** Gold/**NNP** Fields/**NNP** PLC/**NNP** ,/, was/**VBD** named/**VBN** a/**DT** nonexecutive/**JJ** director/**NN** of/**IN** this/**DT** British/**JJ** industrial/**JJ** conglomerate/**NN** ./.

...

38,219 It/**PRP** is/**VBZ** also/**RB** pulling/**VBG** 20/**CD** people/**NNS** out/**IN** of/**IN** Puerto/**NNP** Rico/**NNP** ,/, who/**WP** were/**VBD** helping/**VBG** Hurricane/**NNP** Hugo/**NNP** victims/**NNS** ,/, and/**CC** sending/**VBG** them/**PRP** to/**TO** San/**NNP** Francisco/**NNP** instead/**RB** ./.

- From the training set, induce a function/algorithm that maps new sentences to their tag sequences.

# Overview

- ▶ Recap: The Tagging Problem
- ▶ Log-linear taggers

# Tagging (Sequence Labeling)

- Given a sequence (in NLP, words), assign appropriate labels to each word.
- Many NLP problems can be viewed as sequence labeling:
  - POS Tagging
  - Chunking
  - Named Entity Tagging
- Labels of tokens are dependent on the labels of other tokens in the sequence, particularly their neighbors

Plays well with others.

VBZ RB IN NNS

## Log-Linear Models for Tagging

- ▶ We have an input sentence  $w_{[1:n]} = w_1, w_2, \dots, w_n$   
( $w_i$  is the  $i$ 'th word in the sentence)



## Log-Linear Models for Tagging

- ▶ We have an input sentence  $w_{[1:n]} = w_1, w_2, \dots, w_n$   
( $w_i$  is the  $i$ 'th word in the sentence)
- ▶ We have a tag sequence  $t_{[1:n]} = t_1, t_2, \dots, t_n$   
( $t_i$  is the  $i$ 'th tag in the sentence)

## Log-Linear Models for Tagging

- ▶ We have an input sentence  $w_{[1:n]} = w_1, w_2, \dots, w_n$   
( $w_i$  is the  $i$ 'th word in the sentence)
- ▶ We have a tag sequence  $t_{[1:n]} = t_1, t_2, \dots, t_n$   
( $t_i$  is the  $i$ 'th tag in the sentence)
- ▶ We'll use an log-linear model to define

$$p(t_1, t_2, \dots, t_n | w_1, w_2, \dots, w_n)$$

for any sentence  $w_{[1:n]}$  and tag sequence  $t_{[1:n]}$  of the same length.  
(Note: contrast with HMM that defines  $p(t_1 \dots t_n, w_1 \dots w_n)$ )

## Log-Linear Models for Tagging

- ▶ We have an input sentence  $w_{[1:n]} = w_1, w_2, \dots, w_n$   
( $w_i$  is the  $i$ 'th word in the sentence)
- ▶ We have a tag sequence  $t_{[1:n]} = t_1, t_2, \dots, t_n$   
( $t_i$  is the  $i$ 'th tag in the sentence)
- ▶ We'll use an log-linear model to define

$$p(t_1, t_2, \dots, t_n | w_1, w_2, \dots, w_n)$$

for any sentence  $w_{[1:n]}$  and tag sequence  $t_{[1:n]}$  of the same length.

(Note: contrast with HMM that defines  $p(t_1 \dots t_n, w_1 \dots w_n)$ )

- ▶ Then the most likely tag sequence for  $w_{[1:n]}$  is

$$t_{[1:n]}^* = \operatorname{argmax}_{t_{[1:n]}} p(t_{[1:n]} | w_{[1:n]})$$

How to model  $p(t_{[1:n]}|w_{[1:n]})$ ?

**A Trigram Log-Linear Tagger:**

$$p(t_{[1:n]}|w_{[1:n]}) = \prod_{j=1}^n p(t_j \mid w_1 \dots w_n, t_1 \dots t_{j-1}) \quad \text{Chain rule}$$

How to model  $p(t_{[1:n]}|w_{[1:n]})$ ?

**A Trigram Log-Linear Tagger:**

$$p(t_{[1:n]}|w_{[1:n]}) = \prod_{j=1}^n p(t_j \mid w_1 \dots w_n, t_1 \dots t_{j-1}) \quad \text{Chain rule}$$

$$= \prod_{j=1}^n p(t_j \mid w_1, \dots, w_n, t_{j-2}, t_{j-1})$$

Independence assumptions

- We take  $t_0 = t_{-1} = *$

## How to model $p(t_{[1:n]}|w_{[1:n]})$ ?

### A Trigram Log-Linear Tagger:

$$p(t_{[1:n]}|w_{[1:n]}) = \prod_{j=1}^n p(t_j \mid w_1 \dots w_n, t_1 \dots t_{j-1}) \quad \text{Chain rule}$$

$$= \prod_{j=1}^n p(t_j \mid w_1, \dots, w_n, t_{j-2}, t_{j-1})$$

Independence assumptions

- ▶ We take  $t_0 = t_{-1} = *$
- ▶ Independence assumption: each tag only depends on previous two tags

$$p(t_j|w_1, \dots, w_n, t_1, \dots, t_{j-1}) = p(t_j|w_1, \dots, w_n, t_{j-2}, t_{j-1})$$

# An Example

Hispaniola/**NNP** quickly/**RB** became/**VB** an/**DT** important/**JJ**  
base/**??** from which Spain expanded its empire into the rest of the  
Western Hemisphere .

- There are many possible tags in the position **??**

$\mathcal{Y} = \{NN, NNS, Vt, Vi, IN, DT, \dots\}$

## Representation: Histories

- ▶ A **history** is a 4-tuple  $\langle t_{-2}, t_{-1}, w_{[1:n]}, i \rangle$
  - ▶  $t_{-2}, t_{-1}$  are the previous two tags.
  - ▶  $w_{[1:n]}$  are the  $n$  words in the input sentence.
  - ▶  $i$  is the index of the word being tagged
  - ▶  $\mathcal{X}$  is the set of all possible histories
- 

Hispaniola/**NNP** quickly/**RB** became/**VB** an/**DT** important/**JJ**  
base/**??** from which Spain expanded its empire into the rest of the  
Western Hemisphere .

- ▶  $t_{-2}, t_{-1} = \text{DT, JJ}$
- ▶  $w_{[1:n]} = \langle \text{Hispaniola, quickly, became, ... , Hemisphere, .} \rangle$
- ▶  $i = 6$



## An Example (continued)

- ▶  $\mathcal{X}$  is the set of all possible histories of form  $\langle t_{-2}, t_{-1}, w_{[1:n]}, i \rangle$
  - ▶  $\mathcal{Y} = \{\text{NN}, \text{NNS}, \text{Vt}, \text{Vi}, \text{IN}, \text{DT}, \dots\}$
  - ▶ We have  $m$  features  $f_k : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$  for  $k = 1 \dots m$
- 

For example:

$$f_1(h, t) = \begin{cases} 1 & \text{if current word } w_i \text{ is base and } t = \text{Vt} \\ 0 & \text{otherwise} \end{cases}$$

analogy to  $e(\text{base} | \text{Vt})$  in HMMs

$$f_2(h, t) = \begin{cases} 1 & \text{if current word } w_i \text{ ends in ing and } t = \text{VBG} \\ 0 & \text{otherwise} \end{cases}$$

difficult for HMMs

...

$$f_1(\langle \text{JJ}, \text{DT}, \langle \text{Hispaniola}, \dots \rangle, 6 \rangle, \text{Vt}) = 1$$

$$f_2(\langle \text{JJ}, \text{DT}, \langle \text{Hispaniola}, \dots \rangle, 6 \rangle, \text{Vt}) = 0$$

...

## Training the Log-Linear Model

- ▶ To train a log-linear model, we need a training set  $(x_i, y_i)$  for  $i = 1 \dots n$ . Then search for

$$v^* = \operatorname{argmax}_v \left( \underbrace{\sum_i \log p(y_i | x_i; v)}_{\text{Log-Likelihood}} - \underbrace{\frac{\lambda}{2} \sum_k v_k^2}_{\text{Regularizer}} \right)$$

- ▶ Training set is simply all history/tag pairs seen in the training data

# The Viterbi Algorithm

Problem: for an input  $w_1 \dots w_n$ , find

$$\arg \max_{t_1 \dots t_n} p(t_1 \dots t_n \mid w_1 \dots w_n)$$

We assume that  $p$  takes the form

$$p(t_1 \dots t_n \mid w_1 \dots w_n) = \prod_{i=1}^n q(t_i \mid t_{i-2}, t_{i-1}, w_{[1:n]}, i)$$

(In our case  $q(t_i \mid t_{i-2}, t_{i-1}, w_{[1:n]}, i)$  is the estimate from a log-linear model.)

# The Viterbi Algorithm

- ▶ Define  $n$  to be the length of the sentence
- ▶ Define

$$r(t_1 \dots t_k) = \prod_{i=1}^k q(t_i | t_{i-2}, t_{i-1}, w_{[1:n]}, i)$$

- ▶ Define a dynamic programming table

$\pi(k, u, v)$  = maximum probability of a tag sequence ending  
in tags  $u, v$  at position  $k$

that is,

$$\pi(k, u, v) = \max_{\langle t_1, \dots, t_{k-2} \rangle} r(t_1 \dots t_{k-2}, u, v)$$

## A Recursive Definition

Base case:

$$\pi(0, *, *) = 1$$

**Recursive definition:**

For any  $k \in \{1 \dots n\}$ , for any  $u \in \mathcal{S}_{k-1}$  and  $v \in \mathcal{S}_k$ :

$$\pi(k, u, v) = \max_{t \in \mathcal{S}_{k-2}} \left( \pi(k-1, t, u) \times q(v|t, u, w_{[1:n]}, k) \right)$$

where  $\mathcal{S}_k$  is the set of possible tags at position  $k$

## The Viterbi Algorithm with Backpointers

**Input:** a sentence  $w_1 \dots w_n$ , log-linear model that provides  $q(v|t, u, w_{[1:n]}, i)$  for any tag-trigram  $t, u, v$ , for any  $i \in \{1 \dots n\}$

**Initialization:** Set  $\pi(0, *, *) = 1$ .

**Algorithm:**

- ▶ For  $k = 1 \dots n$ ,
  - ▶ For  $u \in \mathcal{S}_{k-1}, v \in \mathcal{S}_k$ ,

$$\pi(k, u, v) = \max_{t \in \mathcal{S}_{k-2}} (\pi(k-1, t, u) \times q(v|t, u, w_{[1:n]}, k))$$

$$bp(k, u, v) = \arg \max_{t \in \mathcal{S}_{k-2}} (\pi(k-1, t, u) \times q(v|t, u, w_{[1:n]}, k))$$

- ▶ Set  $(t_{n-1}, t_n) = \arg \max_{(u,v)} \pi(n, u, v)$
- ▶ For  $k = (n-2) \dots 1, t_k = bp(k+2, t_{k+1}, t_{k+2})$
- ▶ **Return** the tag sequence  $t_1 \dots t_n$

# FAQ Segmentation: McCallum et. al

- ▶ McCallum et. al compared HMM and log-linear taggers on a *FAQ Segmentation* task
- ▶ Main point: in an HMM, modeling

$$p(word|tag)$$

is difficult in this domain

## FAQ Segmentation: McCallum et. al

<head>X-NNTP-POSTER: NewsHound v1.33

<head>

<head>Archive name: acorn/faq/part2

<head>Frequency: monthly

<head>

<question>2.6) What configuration of serial cable should I use

<answer>

<answer> Here follows a diagram of the necessary connections

<answer>programs to work properly. They are as far as I know t

<answer>agreed upon by commercial comms software developers fo

<answer>

<answer> Pins 1, 4, and 8 must be connected together inside

<answer>is to avoid the well known serial port chip bugs. The



## FAQ Segmentation: Line Features

begins-with-number  
begins-with-ordinal  
begins-with-punctuation  
begins-with-question-word  
begins-with-subject  
blank  
contains-alphanum  
contains-bracketed-number  
contains-http  
contains-non-space  
contains-number  
contains-pipe  
contains-question-mark  
ends-with-question-mark  
first-alpha-is-capitalized  
indented-1-to-4

## FAQ Segmentation: The Log-Linear Tagger

<head>X-NNTP-POSTER: NewsHound v1.33

<head>

<head>Archive name: acorn/faq/part2

<head>Frequency: monthly

<head>

<question>2.6) What configuration of serial cable should I use

Here follows a diagram of the necessary connections

⇒ "tag=question;prev=head;begins-with-number"

"tag=question;prev=head;contains-alphanum"

"tag=question;prev=head;contains-nonspace"

"tag=question;prev=head;contains-number"

"tag=question;prev=head;prev-is-blank"

## FAQ Segmentation: An HMM Tagger

<question>2.6) What configuration of serial cable should I use

- ▶ First solution for  $p(\text{word} \mid \text{tag})$ :

$$\begin{aligned} p(\text{"2.6) What configuration of serial cable should I use"} \mid \text{question}) = \\ e(2.6 \mid \text{question}) \times \\ e(\text{What} \mid \text{question}) \times \\ e(\text{configuration} \mid \text{question}) \times \\ e(\text{of} \mid \text{question}) \times \\ e(\text{serial} \mid \text{question}) \times \\ \dots \end{aligned}$$

- ▶ i.e. have a **language model** for each *tag*

## FAQ Segmentation: McCallum et. al

- ▶ Second solution: first map each sentence to string of features:

`<question>2.6) What configuration of serial cable should I use`

$\Rightarrow$

`<question>begins-with-number contains-alphanum contains-nonspace  
contains-number prev-is-blank`

- ▶ Use a language model again:

$$\begin{aligned} p(\text{"2.6) What configuration of serial cable should I use"} \mid \text{question}) = \\ e(\text{begins-with-number} \mid \text{question}) \times \\ e(\text{contains-alphanum} \mid \text{question}) \times \\ e(\text{contains-nonspace} \mid \text{question}) \times \\ e(\text{contains-number} \mid \text{question}) \times \\ e(\text{prev-is-blank} \mid \text{question}) \times \end{aligned}$$

## FAQ Segmentation: Results

Method	Precision	Recall
ME-Stateless	0.038	0.362
TokenHMM	0.276	0.140
FeatureHMM	0.413	0.529
MEMM	0.867	0.681

- Precision and recall results are for recovering segments

## FAQ Segmentation: Results

Method	Precision	Recall
ME-Stateless	0.038	0.362
TokenHMM	0.276	0.140
FeatureHMM	0.413	0.529
MEMM	0.867	0.681

- ▶ Precision and recall results are for recovering segments
- ▶ ME-stateless is a log-linear model that treats every sentence separately (no dependence between adjacent tags)

## FAQ Segmentation: Results

Method	Precision	Recall
ME-Stateless	0.038	0.362
TokenHMM	0.276	0.140
FeatureHMM	0.413	0.529
MEMM	0.867	0.681

- ▶ Precision and recall results are for recovering segments
- ▶ ME-stateless is a log-linear model that treats every sentence separately (no dependence between adjacent tags)
- ▶ TokenHMM is an HMM with first solution we've just seen

## FAQ Segmentation: Results

Method	Precision	Recall
ME-Stateless	0.038	0.362
TokenHMM	0.276	0.140
FeatureHMM	0.413	0.529
MEMM	0.867	0.681

- ▶ Precision and recall results are for recovering segments
- ▶ ME-stateless is a log-linear model that treats every sentence separately (no dependence between adjacent tags)
- ▶ TokenHMM is an HMM with first solution we've just seen
- ▶ FeatureHMM is an HMM with second solution we've just seen



## FAQ Segmentation: Results

Method	Precision	Recall
ME-Stateless	0.038	0.362
TokenHMM	0.276	0.140
FeatureHMM	0.413	0.529
MEMM	0.867	0.681

- ▶ Precision and recall results are for recovering segments
- ▶ ME-stateless is a log-linear model that treats every sentence separately (no dependence between adjacent tags)
- ▶ TokenHMM is an HMM with first solution we've just seen
- ▶ FeatureHMM is an HMM with second solution we've just seen
- ▶ MEMM is a log-linear trigram tagger (MEMM stands for "Maximum-Entropy Markov Model")

# Summary

- ▶ Key ideas in log-linear taggers:

- ▶ Decompose

$$p(t_1 \dots t_n | w_1 \dots w_n) = \prod_{i=1}^n p(t_i | t_{i-2}, t_{i-1}, w_1 \dots w_n)$$

- ▶ Estimate

$$p(t_i | t_{i-2}, t_{i-1}, w_1 \dots w_n)$$

using a log-linear model

- ▶ For a test sentence  $w_1 \dots w_n$ , use the Viterbi algorithm to find

$$\arg \max_{t_1 \dots t_n} \left( \prod_{i=1}^n p(t_i | t_{i-2}, t_{i-1}, w_1 \dots w_n) \right)$$

- ▶ Key advantage over HMM taggers: flexibility in the features they can use